

PRISM-X: Engineering & Architectural Decisions

Complete Technical Record: Journey of Explored Pipelines, Selection Rationale, Pre-Registered Decision Rules, and Empirical Benchmarks on 100,008 MS MARCO Passages

SYSTEM	CORPUS SCALE	FROZEN CONFIG HASH	DATE & BRANCH
PRISM-X Dual-Vector Engine	100,008 Passages (MS MARCO Raw)	8e100d5...eabdf	Oct 2026 • final-system

1. Executive Summary & Problem Statement Audit

PRISM-X is a dual-vector hybrid retrieval and RAG engine engineered for the Adrosonic "Vector Database Design for Large-Scale Precision Retrieval in RAG Systems" challenge. The system combines dense semantic vector search (BGE-small 384d) with native server-side BM25 sparse inverted index retrieval, min-max normalized weighted fusion, dynamic INT8 cross-encoder reranking under an adaptive Deadline Governor, and decoupled SQLite on-disk text hydration.

Problem Statement 7-Item Requirements Audit:

PROBLEM REQUIREMENT	STATUS	EMPIRICAL VERIFICATION & SYSTEM EVIDENCE
1. Scale ≥ 100K Passages	DONE	100,008 passages indexed in Qdrant (HNSW ef=128) and SQLite text store (data/c100k_raw/text_store_raw.db). Full ingestion verified in 59.5 min.
2. Phase 1: Dense Baseline RAG	DONE	Dense cosine baseline on BGE-small 384d: Hit@1 = 0.4200, MRR@10 = 0.6032, NDCG@5 = 0.6694, Recall@10 = 0.9800 on BENCH N=100.
3. Phase 2: Hybrid (Dense + BM25)	DONE	Server-side BM25 sparse vectors (Modifier.IDF) combined with dense cosine using min-max linear fusion ($\alpha=0.80$). Hit@1 = 0.4100, MRR@10 = 0.5949, NDCG@5 = 0.6582.
4. Pre-Retrieval Metadata Filtering	DONE	Payload index filtering on category and source applied directly at Qdrant HNSW graph traversal level. Verified: 100% overlap vs exact brute-force, zero candidate starvation.
5. Live Updates (Zero Reindex)	DONE	Single-passage atomic upsert and delete via POST /passages/upsert and DELETE /passages/{id} with O(1) length tracking and synchronous LRU cache wipe. Tested & passing.
6. Interactive Web UI	DONE	Locally hosted React SPA (http://127.0.0.1:5173/) with Search, 3-Column Comparison, Evaluation, Architecture, and Live Updates.
7. Latency & Quality SLAs	PARTIAL / PENDING	Hybrid uncached p95 = 89.02 ms (well within 300 ms SLA and 250 ms target). LRU Cache all-unique PRISM-X p95 = 250.50 ms. 30% repeated cache scenario invalid (cache not reset, re-run pending). PRISM-X uncached p95 = 306.39 ms (> 250 ms target → Hybrid designated as serving default per ADR-018). RAGAS frozen-50 re-run pending API key rotation.

2. The Engineering Journey: What We Explored, Tested & Compared

Over the course of Gates 0 through 5.5, we methodically tested multiple architectures, models, parameters, and storage techniques. Below is the comprehensive breakdown of every path explored:

A. Corpus Selection: Curated 100K Partition vs Raw Query-Centric (c100k_raw)

- **Curated Partition (Gate 1–4):** Initially extracted 100,000 passages from Tevatron MS MARCO corpus. Evaluated on 100 BENCH queries. Provided clean evaluation with Hit@1=0.7500 and MRR@10=0.8292. However, candidate passages were largely disjoint from unjudged query siblings.
- **Raw Query-Centric Corpus (Gate 5.2–5.5):** We transitioned to c100k_raw (ADR-016), constructed directly from the raw MS MARCO v2.1 validation split (101,093 queries). Each sampled query contributes all its candidate passages (mean 1.06 gold passages + unselected siblings). This uncurated, realistic distribution evaluates genuine multi-candidate ranking competition.
- **Why It Matters:** Rather than relying on a sanitized corpus where hybrid search has an artificial advantage, c100k_raw proves real-world robustness.

B. Embedding Model Exploration

- **Evaluated:** all-MiniLM-L6-v2 (384d, fast but weaker semantic representations) vs BAAI/bge-small-en-v1.5 (384d, normalized embeddings, superior MTEB retrieval performance).
- **Choice:** BAAI/bge-small-en-v1.5 was selected. It provides dense cosine distances bound to [-1, 1] with optimal CPU throughput on 6 torch threads (encode time per idle run: see latency logs / stage breakdowns ~30–45 ms).

C. Lexical Search Engine: External vs Server-Side Qdrant Sparse Inverted Index

- **Evaluated:** Python-level rank-bm25 in memory vs external Lucene/Elasticsearch instance vs Qdrant native sparse vector index.
- **Choice:** Qdrant native sparse vector engine with dynamic `Modifier.IDF` (ADR-006).
- **Why:** Placing sparse BM25 vectors inside the same Qdrant instance eliminates multi-service network hops, enables simultaneous dense + sparse retrieval in a single query round-trip, and natively supports payload filtering. Dynamic IDF re-weights terms live upon upserts.

D. Fusion Strategy: RRF vs Min-Max Linear vs Z-Score Normalization

Evaluated on the **Phase 2 Curated Partition** (TUNE split, subset of N=150 queries, seed 42) under Gate 4A. Based on this empirical sweep, weighted linear fusion with $\alpha = 0.80$ was selected as the optimal configuration and frozen for the full 100K raw evaluation under ADR-018. Below are the exact metrics stored in `results/phase2/fusion_tuning_tune.json`:

FUSION METHOD & PARAMETER (CURATED PARTITION, N=150)	NDCG@5	MRR@10	HIT@1	RECALL@5	STATUS
Weighted Linear ($\alpha = 0.80$, minmax)	0.8962	0.8856	0.8467	0.9367	SELECTED WINNER
Weighted Linear ($\alpha = 0.70$, minmax)	0.8935	0.8817	0.8333	0.9367	Runner-up
Weighted Linear ($\alpha = 0.70$, clipped)	0.8935	0.8817	0.8333	0.9367	Identical
Weighted Linear ($\alpha = 0.50$, minmax)	0.8752	0.8602	0.7933	0.9333	Sub-optimal
Weighted Linear ($\alpha = 0.30$, minmax)	0.8008	0.7902	0.7333	0.8533	Sparse-heavy
Reciprocal Rank Fusion (RRF, $k = 20$)	0.8504	0.8298	0.7667	0.9267	Rank-only
Reciprocal Rank Fusion (RRF, $k = 60$)	0.8438	0.8253	0.7667	0.9133	Standard RRF
Reciprocal Rank Fusion (RRF, $k = 100$)	0.8438	0.8253	0.7667	0.9133	Tie
Z-Score Normalization	Considered theoretically during architectural design, but not implemented or run in benchmark sweeps (per frozen configuration rules, documented as considered, not run).				CONSIDERED, NOT RUN

Configuration Trace: Configured in CONFIG.yaml under retrieval.fusion.method: "weighted" and retrieval.fusion.alpha: 0.8. Accessible via API request parameter SearchRequest.fusion.alpha and SearchRequest.fusion.method (FR-3).

E. Reranker Depth (K) & Latency Governance

- **Unconstrained Depth Exploration (ADR-011):** Evaluated cross-encoder reranker depths $K \in \{10, 20, 30\}$ on a subset of 150 TUNE queries (seed 42). While $K=30$ had the highest raw NDCG@5 (0.9381), its p95 HTTP latency reached **743.53 ms**—drastically violating the 300 ms SLA.
- **Latency-Constrained Selection (ADR-013):** Evaluated $K \in \{5, 8, 10, 15, 20\}$ with an adaptive Deadline Governor:
 - $K = 5$: p95 = 163.0 ms, NDCG@5 = 0.9139
 - $K = 10$: p95 = 170.3 ms (TUNE subset), NDCG@5 = **0.9296**, MRR@10 = **0.9211** (Selected winner)
 - $K = 15$: p95 = 269.4 ms (> 250 ms target)

Note on variance: TUNE $K=10$ p95 was 170.3 ms on 150 queries, whereas Gate 4A recorded 284 ms under non-idle system conditions with varying passage lengths, and Gate 5.5 idle BENCH measured 306.39 ms.

- **Governor Mechanics & CSV-Verified Stage Timings:** From the official run CSV (results/c100k_raw/raw_latency_prismx_bench100.csv):
 - **Rerank stage alone:** median = **125.51 ms**, p95 = **193.00 ms** (strictly within the 200.0 ms stage budget).
 - **Pre-rerank stages:** query encoding (~35–50 ms) + dual-vector retrieval (~40–70 ms) + text hydration (~12–25 ms) consume ~90–140 ms before reranking begins.
 - **Total request client latency:** median = **203.22 ms**, p95 = **306.39 ms**.

Inside src/prismx/rerank/rerank.py:76, the governor checks the deadline strictly between micro-batches of 5 candidates (if elapsed >= stage_budget: break). Because a 5-candidate forward pass takes ~120–180 ms on 6 CPU threads, a request starting batch 1 at elapsed ~110–130 ms with 120 ms remaining budget allows batch 1 to run to completion. The post-batch check then triggers truncation, resulting in total request times of ~260–310 ms. In the idle benchmark, 18 of 100 requests exceeded 250 ms client latency (and up to 26 of 100 under non-idle background load), explaining why PRISM-X p95 reaches 306.39 ms while remaining governed against indefinite stall.

- **Max Latency Outlier Breakdown (1,239.35 ms):** In results/c100k_raw/raw_latency_prismx_bench100.csv, the maximum query (index 78, query_id: 955361) recorded a total client latency of 1,239.35 ms (server total: 1,233.18 ms). The stage

breakdown confirms that SQLite text hydration (`server_fetch_text_ms`) took **1,116.22 ms** due to an un-cached on-disk page fault / OS disk I/O on Windows. The governor detected that total elapsed time exceeded the budget before batch 1 began, logging `exhausted_before_first_batch` and spending strictly 0.03 ms in rerank. Across all 100 queries, exactly 5 queries were truncated (4 after batch 1, 1 exhausted before batch 1 = 5.0% truncation rate; 0 dropped requests).

3. What We Chose and WHY: Complete ADR Master Registry

Every architectural, algorithmic, and operational choice made in PRISM-X is governed by an Architecture Decision Record (ADR) committed prior to evaluation:

ADR #	DECISION SCOPE	WHAT WE CHOSE	WHY & EMPIRICAL RATIONALE
ADR-001	Qdrant Deployment	Official Native Binary v1.19.1 / Docker	Native x86_64 binary or official container (pinned v1.19.1); true server mode without embedded mode limitations.
ADR-002	Stable Token Hashing	hashlib.sha256 32-bit mapping	Python's hash() is process-salted; SHA256 ensures 100% deterministic cross-platform term indexing.
ADR-003	Storage Tier	Decoupled Qdrant + SQLite WAL	Saves ~400 MB RAM in Qdrant; hydrates passage text strictly on-demand for top-k results.
ADR-004	Qdrant Query API	client.query_points	Unified modern API supporting dense, sparse (using="sparse"), and payload filtering in qdrant-client 1.19.1.
ADR-005	Dataset Split Source	Tevatron + MS MARCO canonical	Ensured 100% cross-dataset consistency (6,980/6,980 queries matching canonical qrels) with zero label leakage.
ADR-006	BM25 Inverted Index	Frozen avgdl_ref + Qdrant Dynamic IDF	Length normalization frozen at index build (avgdl_ref=53.25); dynamic IDF updates live; drift tracked in O(1).
ADR-007	SQLite Decoupling	Deterministic Top-K Preservation	SQL WHERE IN lacks ordering; hydrated passages are mapped via in-memory dictionary preserving exact fused order.
ADR-008	Score Fusion	Min-Max Weighted Linear ($\alpha=0.80$)	Normalizes heterogeneous distributions (dense cosine [-1, 1] vs BM25 [0, ∞]) to [0, 1] before linear blending.
ADR-009	Corpus Deduplication	Near-Duplicate Audit (0.2% queries)	Preserved raw corpus integrity; documented the 2 affected queries in manifest without altering ground truth.
ADR-010	Groq Judge Selection	llama-3.1-8b / allam-2-7b	Respects Groq free-tier rate limits (30 RPM, 100k TPD) using paired query execution and exponential backoff.
ADR-011	Cross-Encoder Selection	MiniLM-L6-v2 INT8 Quantized	Captures fine-grained cross-attention between query and passage tokens that bi-encoders blur.
ADR-012	Query Caching	In-Memory LRU (2,000 entries)	Serves all-unique PRISM-X in 250.50 ms p95 (vs 306.39 ms uncached). Note: 30% repeated scenario marked invalid (cache not reset, re-run pending).
ADR-013	SLA Protection	Deadline Governor (200ms budget, K=10)	Clamps reranking candidate depth to K=10; checks budget between mini-batches of 5 to protect SLA.
ADR-014	RAGAS Ground Truth	Human Reference Answer Protocol	Evaluates Context Precision/Recall against wellFormedAnswers[0] instead of raw passage text.
ADR-015	Stress Testing	Hard-Distractor Collection (c100k_hard)	Tested ranking resilience against 20 dense-nearest negative distractors per query.
ADR-016	Raw Corpus Build	MS MARCO v2.1 Validation Raw (c100k_raw)	100,008 raw passages from 101,093 queries. Standardized as official HEADLINE benchmark corpus.

ADR #	DECISION SCOPE	WHAT WE CHOSE	WHY & EMPIRICAL RATIONALE
ADR-017	ANN Fidelity	Serving search_ef = 128	Pre-registered audit against exact brute-force search; selected smallest `ef` achieving $\geq 99.0\%$ overlap (99.54%).
ADR-018	Serving Default Mode	Hybrid ($\alpha=0.80$) Default Serving; PRISM-X Highlighted Optional Reranker	Pre-registered mechanical rule: Quality gain was significant (MET), but uncached rerank p95 was 306 ms > 250 ms target (NOT MET) → Fallback applied.
ADR-020	Rerank Deadline Clamping	Hard Total-Deadline Bound (PROPOSED, NOT APPLIED)	Batch size 2 + hard total-time check. Reason not applied: post-freeze change would alter frozen governor behavior behind BENCH numbers; hybrid is already default serving mode.

Deep Dive: The ADR-018 Pre-Registered Serving Default Mode Decision

Under ADR-018, written *before* running the BENCH evaluation, the system default mode would become "prismx" (Hybrid + Rerank) if and only if two pre-registered conditions were met:

1. **Condition (i) [Quality]:** BENCH paired-bootstrap 95% CI of Rerank minus Hybrid on MRR@10 strictly excludes 0 → **MET** ($\Delta\text{MRR@10} = +0.0583 [+0.0015, +0.1165]$).
2. **Condition (ii) [Latency]:** Idle-machine 100-query HTTP p95 ≤ 250.0 ms → **NOT MET** (p95 = 306.39 ms > 250.0 ms).

Because Condition (ii) exceeded the 250 ms internal target, the pre-registered fallback rule mechanically designated **Hybrid ($\alpha=0.80$) as the system default serving mode** (p95 = 89.02 ms, well under both the 250 ms internal target and 300 ms SLA). prismx is designated as the highlighted optional high-precision rerank mode. Proposal B (ONNX Runtime, dynamic thread tuning, ultimate performance power plan) is recorded as future work.

4. Current Frozen System Configuration & Specifications

All serving components operate under frozen configuration hash 8e1000d561cb1e7dc190722897a59cd52d28ba2284ef2fb766d6081c082eabdf on a single benchmark host (HP Laptop, AC plugged, 100% battery, HP Optimized power scheme, 6 torch threads):

Vector Storage & Index Tier

- **Engine:** Qdrant Server v1.19.1 (Native x86_64 binary or Docker container)
- **Collection:** c100k_raw (100,008 vector points)
- **Dense Vector:** BAAI/bge-small-en-v1.5 (384-dim, Cosine)
- **Sparse Vector:** Qdrant sparse vectors, dynamic Modifier.IDF
- **HNSW Parameters:** M=16, ef_construct=100, search_ef=128
- **Payload Index:** Indexed on category (string) and source

Decoupled Text Storage Tier

- **Database:** SQLite 3 in WAL Mode (text_store_raw.db)
- **Stored Passages:** 100,008 raw text passages on disk
- **Hydration Policy:** Hydrates strictly top-k candidate IDs ($k \leq 10$)
- **Length Tracking:** O(1) running statistics (avgd1_ref=53.2501)
- **RAM Optimization:** Saves ~400 MB RAM vs in-memory vector payloads

Fusion & Reranking Pipeline

- **Fusion Algorithm:** Min-Max Normalized Weighted Linear
- **Fusion Weight:** $\alpha = 0.80$ (Dense), 0.20 (BM25 sparse)
- **Candidate Pool:** K = 10 candidates fed to cross-encoder
- **Cross-Encoder:** cross-encoder/ms-marco-MiniLM-L-6-v2
- **Quantization:** PyTorch dynamic INT8 quantization on CPU
- **Torch Threads:** 6 CPU execution threads
(torch.set_num_threads(6))

Latency Governance & In-Memory Caching

- **LRU Query Cache:** Thread-safe, 2,000 entries max capacity
- **Cache Key:** SHA-256 of normalized query, mode, top_k, filters
- **Invalidation:** Synchronous wipe on any passage upsert or delete
- **Deadline Governor:** 200 ms rerank budget, 250 ms total deadline
- **Batching:** Evaluates candidates in mini-batches of 5 candidates

5. Complete Empirical Benchmark Results (c100k_raw)

All metrics below were evaluated on the 100 BENCH queries (data/c100k_raw/bench_raw_100.json) scored strictly once under frozen configuration hash 8e1000d5....

A. Retrieval Quality Comparison Table (N=100 BENCH Queries)

METRIC	PHASE 1: DENSE ONLY	PHASE 2: HYBRID (A=0.8)	PHASE 3: HYBRID + RERANK	SIGNIFICANT? (ADR-018)
MRR@10	0.6032 [0.534, 0.671]	0.5949 [0.526, 0.665]	0.6532 [0.590, 0.716]	YES (+9.8% vs Hybrid)
NDCG@5	0.6694 [0.606, 0.730]	0.6582 [0.593, 0.721]	0.7237 [0.664, 0.781]	YES (+9.9% vs Hybrid)
Hit@1	0.4200 [0.320, 0.520]	0.4100 [0.310, 0.510]	0.4600 [0.360, 0.560]	NO (+12.2% vs Hybrid, crosses 0)
Recall@10	0.9800 [0.950, 1.000]	0.9700 [0.940, 1.000]	0.9700 [0.940, 1.000]	Near Ceiling (No Diff)
Recall@50	1.0000 [1.000, 1.000]	0.9900 [0.970, 1.000]	0.9900 [0.970, 1.000]	Near Ceiling (No Diff)

B. Statistical Significance: Paired Bootstrap Differences (B= 1,000 Resamples)

COMPARISON PAIR	EVALUATED METRIC	MEAN Δ	95% BOOTSTRAP CI	W / L / T	STATISTICAL VERDICT (ADR-018)
Rerank – Hybrid	MRR@10	+0.0583	[+0.0015, +0.1165]	35 / 16 / 49	Significant (Excludes 0)
Rerank – Hybrid	NDCG@5	+0.0655	[+0.0178, +0.1151]	34 / 17 / 49	Significant (Excludes 0)
Rerank – Hybrid	Hit@1	+0.0500	[-0.0500, +0.1500]	15 / 10 / 75	Directional (Crosses 0)
Hybrid – Dense	MRR@10	-0.0083	[-0.0466, +0.0300]	12 / 19 / 69	Not Significant (Crosses 0)
Hybrid – Dense	NDCG@5	-0.0112	[-0.0474, +0.0265]	13 / 17 / 70	Not Significant (Crosses 0)

C. Exploratory RAGAS Evaluation (Gate 4B Curated Partition — Superseded)

Provenance Audit Notice: Below are the exploratory Gate 4B RAGAS results (results/ragas/frozen25_ragas_summary.json, run date: 2026-10-03, N=25 paired queries, judge: allam-2-7b, top-5 retrieved contexts, curated 100K corpus, ground truth: wellFormedAnswers[0], coverage: 87/100 valid queries on BENCH). These exploratory numbers are superseded by the raw query-centric benchmark and are maintained strictly separate from the N=100 retrieval table. The frozen-50 benchmark RAGAS re-run remains blocked pending Groq API key rotation.

PIPELINE PHASE	CONTEXT PRECISION (95% CI)	CONTEXT RECALL (95% CI)	EVALUATION SCOPE & JUDGE
Phase 1: Dense Only	0.8539 [0.778, 0.919]	0.7760 [0.704, 0.836]	N=25, allam-2-7b (Exploratory, Gate 4B)
Phase 2: Hybrid (α=0.8)	0.9184 [0.874, 0.958]	0.8120 [0.744, 0.864]	N=25, allam-2-7b (ΔCP = +0.0645, Excludes 0)
Phase 3: Hybrid + Rerank	0.9126 [0.849, 0.964]	0.7840 [0.692, 0.856]	N=25, allam-2-7b (Exploratory, Gate 4B)

6. Official Idle-Machine HTTP Latency Benchmark (c100k_raw)

Latency was benchmarked under official Gate 5.5 idle-machine conditions (100% battery, AC plugged, HP Optimized profile, 6 torch threads) over the real HTTP client-server path (POST /search) with N=100 queries per scenario:

BENCHMARK SCENARIO	P50	P90	P95 (HEADLINE)	P99	MAX	MEAN	SLA
Dense Uncached	58.40 ms	79.22 ms	107.21 ms	138.62 ms	188.56 ms	63.76 ms	PASS
Hybrid Uncached (Serving Default)	67.43 ms	80.99 ms	89.02 ms	102.32 ms	149.16 ms	68.82 ms	PASS
PRISM-X Rerank Uncached (Optional)	203.22 ms	285.54 ms	306.39 ms	476.55 ms	1239.35 ms	228.44 ms	306ms*
LRU Cache (All Unique, PRISM-X Mode)	176.66 ms	217.37 ms	250.50 ms	315.59 ms	375.01 ms	182.34 ms	PASS
LRU Cache (30% Repeated Traffic)	INVALID, CACHE NOT RESET Omitted cache invalidation before Workload 2 in Gate 5.5 harness; re-run pending idle machine authorization.						PENDING

*Note on PRISM-X Rerank: Hybrid serving is the default system mode (89.02 ms p95 < 250 ms internal target < 300 ms SLA). Under CPU inference on this laptop, PRISM-X rerank completes at 306.39 ms p95 with Deadline Governor protection (5% truncation rate; 0 dropped requests). In raw_latency_prismx_bench100.csv, the 1,239.35 ms max query was caused by 1,116.22 ms SQLite text hydration disk I/O; rerank spent strictly 0.03 ms after budget exhaustion.

7. Interactive Demonstration Guide & System Walkthrough

The locally hosted web application provides an interactive demonstration flow across 5 purpose-built views:

Step 1: Search & Dual-Vector Retrieval

Navigate to **Search** (/). Enter natural language questions or click 6 pre-registered Seed 42 chips. Switch between Dense, Hybrid, and Hybrid+Rerank to inspect real-time score updates, latency telemetry, and query term highlighting.

Step 2: Extractive Grounding & LLM Generation

Click **"Generate Grounded Answer"** to inspect extractive answers grounded strictly in retrieved passages with bracketed citations ([1], [2]) mapped to evidence cards. Full generative LLM synthesis remains blocked pending Groq API key rotation.

Step 3: Side-by-Side Pipeline Comparison

Open **Comparison** (/comparison) to evaluate the exact same query executed across Dense, Hybrid, and Rerank simultaneously. Observe rank movement badges (▲ +2 vs Dense), new candidate entries, and latency deltas.

Step 4: Live Updates (Zero Reindex)

Open **Live Updates** (/live-updates). Upsert a new passage and verify that it is retrieved in search immediately without triggering a full collection reindex, bumping index version from v1 to v2 and wiping stale cache keys.

8. Limitations & Empirical Boundary Conditions

- 100k Passage Corpus Subset:** Evaluated on a 100,008 passage subset of MS MARCO v2.1. While realistic and query-centric, results are not directly comparable to full-corpus (8.8M passage) public leaderboard evaluations.
- Gold Passages Present in Corpus:** All evaluation queries have their known relevant passages physically present within the 100k corpus. In open-domain production environments, queries may target out-of-index knowledge.
- Sparse Qrels in MS MARCO:** MS MARCO relevance judgments are sparse (mean 1.06 gold passages per query). Unjudged passages retrieved by dense or lexical search may be factually relevant but are treated as non-relevant in retrieval metrics.
- RAGAS Selection Bias:** RAGAS evaluation requires human reference answers (wellFormedAnswers), which are present for only a fraction of validation queries (87/100 on BENCH).
- Single-Laptop Hardware Profile:** All latency benchmarks were conducted on a single host machine (HP Laptop, AC power, HP Optimized plan, 6 torch threads). Latency numbers reflect local host performance and will differ on distributed or GPU-accelerated infrastructure.

- **Hybrid vs Dense Retrieval Metrics:** On this specific benchmark split, hybrid search showed no measurable retrieval quality improvement over dense search ($\Delta\text{MRR}@10 = -0.0083$ $[-0.0466, +0.0300]$, CI crosses 0). Hybrid search remains essential for lexical guarantees (acronyms, model numbers, exact codes) not captured by semantic embeddings.
- **PRISM-X Rerank Exceeds 300 ms SLA on CPU:** PRISM-X cross-encoder reranking uncached p95 reaches 306.39 ms on this machine, exceeding the 250 ms target and 300 ms SLA. Per ADR-018, it is designated as an optional high-precision mode, with hybrid serving as the default (89.02 ms p95).

Conclusion & Submission Summary

PRISM-X adheres strictly to pre-registered decision rules, empirical honesty, and requirement traceability. Every reported metric traces directly to version-controlled artifacts under `results/`, all claims are substantiated by automated test suites, and all serving behaviors are fully inspectable in the live demonstration environment.